

Image_convolution

January 14, 2016

```
In [1]: # This is a cell to hide code snippets from displaying
        # This must be at first cell!
```

```
from IPython.display import HTML, Image
```

```
hide_me = ''
HTML('''<script>
code_show=true;
function code_toggle() {
    if (code_show) {
        $('div.input').each(function(id) {
            el = $(this).find('.cm-variable:first');
            if (id == 0 || el.text() == 'hide_me') {
                $(this).hide();
            }
        });
        $('div.output_prompt').css('opacity', 0);
    } else {
        $('div.input').each(function(id) {
            $(this).show();
        });
        $('div.output_prompt').css('opacity', 1);
    }
    code_show = !code_show
}
$( document ).ready(code_toggle);
</script>
<form action="javascript:code_toggle()"><input style="opacity:0" type="submit" value="Click here to toggle the code" /></form>''')
```

```
Out[1]: <IPython.core.display.HTML at 0x7f945c47aa50>
```

1 Image convolution: Direct space versus Fourier space

In image processing, the convolution with a small matrix is used for blurring, sharpening, embossing, edge-detection, and more. This is accomplished by means of convolution between a kernel and an image. http://en.wikipedia.org/wiki/Kernel_%28image_processing%29

In this exercise we will focus on the gaussian blur and compare the execution time of two approaches: direct space vs Fourier space for different width of gaussians: 1, 2, 4, 8, and 16pixels.

The raw image used is the famous 512x512 pixel image from Lena, and available as part of scipy.

Nota: Since beginning of 2016, the image of Lena is planned for removal from scipy. One should use `scipy.misc.ascent()` instead of `scipy.misc.lena()`


1.1 Loading the Scientific Python stack and the input image

```
In [2]: %pylab inline
```

Populating the interactive namespace from numpy and matplotlib

```
In [3]: from scipy.misc import lena
        img = lena()
        imshow(img, cmap="gray")
```

```
Out[3]: <matplotlib.image.AxesImage at 0x7f943f422910>
```



1.2 Properties of the gaussian for blurring an image

```
In [4]: hide_me
        from IPython.display import Image
        Image(filename="files/gaussian.png")
```


```
Out[4]:
```

$$g(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}.$$

The gaussian has some key advantages: * Very smooth blurring * Separable in x/y (hence much faster) * The Fourier transform is a gaussian of width proportional to 1/sigma

```
In [5]: from scipy.signal import gaussian
        plot(gaussian(21,1))
```

```
Out[5]: [<matplotlib.lines.Line2D at 0x7f94340c2a10>]
```



1.3 Direct space approach

Choice of the window size ?

One of the key parameter of a direct space convolution is the size of the window for the convolution. To calculate it, we take a gaussian with $\sigma = 1.0$ and evaluate the number of integer position have non-neglectable values (>0.0001).

```
In [6]: (gaussian(21,1)>1e-4).sum()
```

```
Out[6]: 9
```

It is important to have an odd window size, we choose $w = 8 \sigma + 1$.


scipy.signal has a special function for direct space convolution of separable functions: *sepfir2d*
<https://docs.scipy.org/doc/scipy-0.7.x/reference/generated/scipy.signal.sepfir2d.html>

Performing 2 times a 1D convolution is much faster than performing a single 2D convolution.

```
In [7]: from scipy.signal import sepfir2d
def direct(img, sigma):
    gaus = gaussian(8 * sigma + 1, sigma)
    return sepfir2d(img, gaus, gaus)
```

```
In [8]: imshow(direct(img, 5), cmap="gray")
```

```
Out[8]: <matplotlib.image.AxesImage at 0x7f9434063d50>
```



```
In [9]: %timeit direct(img, 2)
```

100 loops, best of 3: 12.4 ms per loop

1.4 Fourier space convolution

The convolution in direct space is equivalent to a multiplication in Fourier space.
http://en.wikipedia.org/wiki/Convolution_theorem

```
In [10]: hide_me
         from IPython.display import Image
         Image(filename="files/convolution.png")
```

Out[10]:


$$f * g = \mathcal{F}^{-1} \{ \mathcal{F}\{f\} \cdot \mathcal{F}\{g\} \}$$

This is especially appealing for gaussian blurring as the FT of a gaussian is a gaussian. The total cost is reduced to one FT and its inverse.

```
In [11]: from scipy.fftpack import fft2, ifft2, fftshift
         def ft(img, sigma):
             size = len(img)
             gaus = gaussian(size, size / (2*pi*sigma))
             sgau = fftshift(gaus)
             return ifft2(fft2(img) * outer(sgau,sgau)).real
```

```
In [12]: imshow(ft(img, 5), cmap="gray")
```

Out[12]: <matplotlib.image.AxesImage at 0x7f942f8451d0>



```
In [13]: %timeit ft(img, 2)
```

10 loops, best of 3: 26.8 ms per loop


1.5 Benchmark: Direct space vs Fourier space

```
In [14]: import time
         sigmas = []
         d = []
         f = []
         for i in range(5):
             sigma = 1<<i
             t0 = time.time()
             direct(img, sigma)
             t1 = time.time()
             ft(img, sigma)
             t2 = time.time()

             print("sigma=%s \t Direct: %.3fs \t Fourier: %.3fs"%(sigma,t1-t0,t2-t1))
             sigmas.append(sigma)
             d.append(t1-t0)
             f.append(t2-t1)
         plot(sigmas,d,label="Direct")
         plot(sigmas,f,label="Fourier")
         legend()
```

sigma=1	Direct: 0.008s	Fourier: 0.029s
sigma=2	Direct: 0.011s	Fourier: 0.026s
sigma=4	Direct: 0.021s	Fourier: 0.026s
sigma=8	Direct: 0.042s	Fourier: 0.026s
sigma=16	Direct: 0.089s	Fourier: 0.029s

Out[14]: <matplotlib.legend.Legend at 0x7f942f78b650>



Nota: The gaussian blur filter exists also as “ready to use” function as part of `scipy.ndimage.filters.gaussian_filter`. We can compare how it behaves compared to our implementations:

```
In [15]: from scipy.ndimage import filters
```

```
In [16]: %timeit filters.gaussian_filter(img, 2)
```

100 loops, best of 3: 7.1 ms per loop

```
In [17]: %timeit filters.gaussian_filter(img, 4)
```

100 loops, best of 3: 10.5 ms per loop

```
In [18]: import time
         from scipy.ndimage import filters
         sigmas = []
         d = []
         f = []
         n = []
         for i in range(5):
             sigma = 1<<i
             t0 = time.time()
             direct(img, sigma)
```


```

t1 = time.time()
ft(img, sigma)
t2 = time.time()
filters.gaussian_filter(img, sigma)
t3=time.time()
print("sigma=%s \t Direct: %.3fs \t Fourier: %.3fs, \t ndimage: %.3fs"%(sigma,t1-t0,t2-t1,
sigmas.append(sigma)
d.append(t1-t0)
f.append(t2-t1)
n.append(t3-t2)
plot(sigmas,d,label="Direct")
plot(sigmas,f,label="Fourier")
plot(sigmas,n,label="Ndimage")
legend()

```

sigma=1	Direct: 0.007s	Fourier: 0.041s,	ndimage: 0.007s
sigma=2	Direct: 0.011s	Fourier: 0.026s,	ndimage: 0.008s
sigma=4	Direct: 0.021s	Fourier: 0.026s,	ndimage: 0.011s
sigma=8	Direct: 0.047s	Fourier: 0.026s,	ndimage: 0.019s
sigma=16	Direct: 0.090s	Fourier: 0.029s,	ndimage: 0.037s

Out[18]: <matplotlib.legend.Legend at 0x7f942f66f910>



2 Conclusion

The choice of the algorithm is of prime importance for image manipulation: Fourier transform are extremely fast and optimized but they cannot compete with simpler algorithms when the window is small. Moreover, if your problem is common, there is probably an optimized implementation already available within *scipy*.